A1: Games

Deadline

For grade **A** you must have your implementations accepted by Kattis with the required scores by **Feb 7 2020, 23:59 sharp** and present in one of the "on-time" slots that will be scheduled during the weeks Feb 24-28 and March 2-6. Any assignment accepted by Kattis/presented after these deadlines will result in max. grade B.

1 INTRODUCTION TO GAME THEORY

Game theory studies systems where two or more agents try to maximize their own gain by operating on a shared stated. In this assignment, we are interested in a special class of such problems, namely one where two players are in direct conflict (which are commonly called zero-sum games).

To put it formally, a game consists of the following parts:

$P = \{A, B\}$ is the list of players (which will always be the same in our case because we will only analyze two-player games).

S is the list of possible game states.

s_0 is the initial state, the game *always starts* in the initial state.

$\mu : P \times S \rightarrow \mathscr{P}(S)$ (where $\mathscr{P}(S)$ is the set of all subsets of S) is a function that given a player and a game state returns the possible game states that the player may achieve with one *legal* move by the current player.

$\gamma : P \times S \rightarrow \mathbb{R}$ is a *utility function* that given a player and a state says how "useful" the state is for the player.

A game is called a zero-sum game if $\gamma(A, s) + \gamma(B, s) = 0$ or equivalently $\gamma(A, s) = -\gamma(B, s)$.

Informally, each component of the specification has a role in structuring the gameplay. More concretely we have:

Figure 1.1: Initial game state for checkers.

Not all states in S need to be reachable by a legal sequence of moves.

μ is the part of the specification that dictates the *rules* of the game (by dictating which are the legal moves).

γ is the part of the specification that drives the game forward, and it is not unique (any rescaling of γ is also a utility function).

We say that the game is over when it reaches a state s_t such that $\mu(p, s) = \emptyset$ and player p is about to play. In this case s is called a *terminal state*, and we say that:

Player A wins if $\gamma(A, s) > \gamma(B, s)$.

Player B wins if $\gamma(B, s) > \gamma(A, s)$.

The game is tied if $\gamma(A, s) = \gamma(B, s) = 0$.

Hence the objective of player A is to reach a terminal state s_t that maximizes $\gamma(A, s_t)$ and since $\gamma(B, s) = -\gamma(A, s)$, the objective of player B is to reach a terminal state s_t' that minimizes $\gamma(A, s_t')$.

EXAMPLE: CHECKERS

- S is any 8×8 checker board where only the black squares are occupied, and there are at most 12 checkers of each color.

- s_0 is the board with 12 checkers of each color positioned in opposite sides of the board, as shown in figure 1.1.

- μ assigns to a game state (s, A) the set $\mu(s, A)$ comprising the states s' that can be obtained from s by moving a white checker (or a red checker in case of B).

Describe the possible states, initial state, transition function.

> Describe the terminal states of both checkers and tic-tac-toe.

1.1 HEURISTIC FUNCTIONS

Note that in the previous example we did not describe any utility function. This is because utility functions are used as a theoretical device and intuitively they describe how a perfect player would play (at each step the perfect player would choose the next state with the highest utility). It is natural to assume that such a function is difficult to obtain, so instead one generally uses *heuristic functions.*

Simply put, a heuristic function is one that approximates a utility function. Usually the process for obtaining such a function is by analyzing a simpler problem, so that the resulting function is easier to compute. With that said, a rather simple heuristic function for checkers could be given by:

$$v(A, s) = \#\{\text{white checkers}\} - \#\{\text{red checkers}\}$$

This function v belongs to a class of heuristic functions called *evaluation functions* meaning that it can be efficiently computed using only the game state[1].

> Why is $v(A, s) = \#\{\text{white checkers}\} - \#\{\text{red checkers}\}$ a valid heuristic function for checkers (knowing that A plays white and B plays red)?

> When does v best approximate the utility function, and why?

> Can you provide an example of a state s where $v(A, s) > 0$ and B wins in the following turn? (Hint: recall the rules for jumping in checkers)

1.2 GAME TREES AND MORE HEURISTICS

From the questions in the previous section (specifically the last one) it should be evident that a naive heuristic function is, in general not enough to guarantee a victory against a smart (and sometimes, even random) player. This shortcoming is shared by most evaluation functions. In this section we introduce the minimax algorithm that, given a naive heuristic function produces a heuristic function that better fits a utility function. To this end we start by introducing the following heuristic function:

$$\eta(A, s) = |w(s, A)| - |l(s, A)| \qquad \eta(B, s) = |w(s, B)| - |l(s, B)|$$

---

[1] Note that given any heuristic function it would be, theoretically possible to create a hash-table associating each possible state with its utility function, and this would constitute an evaluation function.

Figure 1.2: Each node corresponds to a game state, The blue nodes have even depth, the red nodes have odd depth, and there is an edge from a blue node to a red node (with corresponding states s_b, s_r, respectively) if in state s_b, player A can perform a move that will change the state into s_r and there is an edge from a red node to a blue node (with corresponding states s'_r, s'_b, respectively) if in state s'_r, player B can perform a move that transforms state into s'_b.

Where $w(s, p)$ is the set of terminal states at which player p wins and which are reachable from state s, and $l(s, p)$ is the set of terminal states at which player p loses and which are reachable from state s (and $|\cdot|$ denotes the list length).

It is easy to see that η satisfies $\eta(A, s) + \eta(B, s) = 0$, and for a terminal state s_t $\eta(A, s_t) > 0$ if player A wins, $\eta(A, s_t) < 0$ if player B wins, and $\eta(A, s_t) = 0$ if it is a draw. Also, note that η is not an evaluation function, as it's computation requires one to filter through valid sequences of moves.

Now, for simplicity assume that player A always starts the game. In order to calculate $\eta(A, s)$ it is convenient to represent the game as a *game tree*, $GT = (V, E)$ with a set of nodes V, each of which corresponds to a game state[2] and a set of edges $E \subseteq V \times V$ satisfying: *Given $s, s' \in V$, the edge (s, s') (also denoted $s \to s'$) is in E if and only if $s' \in \mu(A, s)$ and depth(s) is even, or $s' \in \mu(B, s)$ and depth(s) is odd (see illustration in figure 1.2).*

A game tree can be used to search for a winning strategy at each point of the game. Particularly, the *complete game tree* is a game tree with s_0 as root and where every possible move is expanded until it reaches a terminal state. We can use game trees to solve games, i.e., to find a sequence of moves that one of the players can follow to guarantee a win. Thus, if we have access to the complete game tree, we can calculate $\eta(A, s)$ by checking all terminal states that are reachable from s by a path of strictly increasing depth.

Will η suffer from the same problem (referred to in the last question) as the evaluation function v?

[2]Note that there may be more than one node with the same state.

Figure 1.3: Consider this chessboard setup, with black to play. While playing white an algorithm that only uses a heuristic function such as η may treat the black queen moving horizontally to $a4$ in the same way as moving diagonally to $a6$ while any player would be able to see that moving the queen to $a4$ would result in the loss of the queen, and most likely the loss of any chance of mate by the black player, whereas moving the queen diagonally would result in mate in two turns.

1.3 MINIMAX ALGORITHM

Now even though η represents a step in the right direction as it takes information about how the game might proceed from every point, it still has a large drawback. It assumes player B plays randomly, and assigns the same weight to moves of player B independent of their outcome (see Figure 1.3 for an example).

In the specific case of η, consider for instance $\eta(A, s) = 10$ for a given state s. This just means that there are 10 more reachable terminal states that result in a victory for player A than for player B, however it tells nothing of the distribution of such states, and how much B can do to counteract those moves.

The way around this problem is assuming that the opponent is also aware that it has to optimize it's chances of winning, and therefore when calculating a heuristic function that requires traversing the tree, we should only consider the transitions that maximize the heuristic function for player B (therefore minimizing the heuristic for player A). This can be done by employing the *minimax algorithm*, the pseudocode is written in Algorithm 1.

Algorithm 1: Pseudocode for minimax algorithm

```
1    int minimax (state, player)
2        // state: the current state we are analyzing
3        // player: the current player
4        // returns a heuristic value that approximates a utility function of the state
5
6        if μ(state, player) = ∅  // terminal state
7            return γ(player, state)
8            // γ is an evaluation function
9
10       else // can search deeper
11
12           if player = A
13               bestPossible = −∞
14               for each child in μ(A, state)
15                   v = minimax(child, B)
16                   bestPossible = max(bestPossible, v)
17               return bestPossible
18
19           else // player = B
20               bestPossible = +∞
21               for each child in μ(B, state)
22                   v = minimax(child, A)
23                   bestPossible = min(bestPossible, v)
24               return bestPossible
```

Note that in Algorithm 1 requires one to evaluate the tree until it reaches a terminal state. However, given any evaluation function which can assign a heuristic for the game, we can truncate the search at a given depth by including the number of levels left to analyze as an argument and decreasing it in each step.

1.4 ALPHA-BETA PRUNING

Now, Algorithm 1 requires one to search the complete game tree (at least until a given depth). However, this may not be the case. The idea behind the *alpha-beta pruning* algorithm is that in some situations, parts of the tree may be ignored if we know a priori that no better result may be achieved.

To understand how this is possible, consider that player A is traversing the tree in order to find the next best move using the minimax algorithm. Instead of just keeping track of the best node found so far, the player can also keep track of the best heuristic value computed so far, α as well as the worst such value β (which yields the best result for B). At each transition in the tree α should be updated, if it is player A's turn, and otherwise β should be updated. Finally, the remainder of a branch should be disregarded whenever α is greater then β since that indicates the presence of a non-desirable state.

Algorithm 2: Pseudocode for minimax algorithm with alpha-beta pruning

```
1    int alphabeta (state, depth, α, β, player)
2        // state: the current state we are analyzing
3        // α: the current best value achievable by A
4        // β: the current best value achievable by B
5        // player: the current player
6        // returns the minimax value of the state
7
8        if depth = 0 or μ(state, player) = ∅
9            // terminal state
10           v = γ(player, state)
11
12       elseif player = A
13           v = −∞
14           for each child in μ(A, state)
15               v = max( v, alphabeta(child, depth-1, α, β, B) )
16               α = max(α, v)
17               if β ≤ α
18                   break // β prune
19
20       else // player = B
21           v = ∞
22           for each child in μ(B, state)
23               v = min(v, alphabeta(child, depth-1, α, β, A))
24               β = min(β, v)
25               if β ≤ α
26                   break // α prune
27       return v
```

2 ASSIGNMENT: N-DIMENSIONAL TIC-TAC-TOE

In this assignment we consider a special case of n-dimensional generalization of Tic-Tac-Toe game. The rules are simple. There is an n-dimensional hypercube H, consisting of 4^n cells. Two players, A and B, take turns marking blank cells in H. The first player to mark 4 cells along a row wins. Here by a row we mean any 4 cells, whose centers lie along a straight line in H. For instance, in the case $n = 2$ any horizontal, vertical and diagonal row is winning. In the case $n = 3$, winning rows lie along the 48 orthogonal rows (those which are parallel to one of the edges of the cube), the 24 diagonal rows, or the 4 main diagonals of the cube, making 76 winning rows in total. Player A always starts the game.

In this assignment, our goal is to find the best possible move for player X given a particular state of the game. The number n and the maximal execution time of the program depend on the grade level.

3 ASSIGNMENT: CHECKERS

For obtaining a higher grade, the student is expected to be able to use the general framework developed for tic-tac-toe and adapt it to play checkers.

Note, that this should amount to simply designing a new heuristic function that operates on a checker board rather than a tic-tac-toe board.

4 GRADE LEVEL E AND D

In order to get E or D, one has to write a program computing the best possible next move for player X, given a particular state of the game for $n = 2$. In this case, it is enough to implement Minimax algorithm on a complete game tree, and speed up the search using the **minimax** algorithm with **alpha-beta pruning**. Due to time constraints, one cannot analyse the complete game tree. Therefore the student should investigate th influence of he maximum search depth and come up with a suitable evaluation function. The total execution time for 100 different board states should not exceed 1s. As an example of a simple evaluation function, one can consider the following.

$$\text{eval(Board)} = \sum_{r \in \text{Rows}} \text{myMarks}(r) + \sum_{c \in \text{Columns}} \text{myMarks}(c) + \sum_{d \in \text{Diagonals}} \text{myMarks}(d)$$

The problem can be found on Kattis:

$$\texttt{https://kth.kattis.com/problems/kth.ai.tictactoe2d}$$

To achieve a passing grade the student has to obtain, at least, 22 out of 25 points on Kattis. Note: The Kattis score does not guarantee the student is given the corresponding grade level, the grade level will depend on later oral examination.

5 GRADE LEVEL C

In order to get C, one has to write a program computing the best possible next move for player X, given a particular state of a game of 3D tic-tac-toe. Like in the previous grade level, due to time constraints, the total execution time for analyzing 100 different board states should not exceed 6s. Just as before, the student should investigate the influence of search depth in the execution time and performance of the algorithm.

https://kth.kattis.com/problems/kth.ai.tictactoe3d

To achieve this grade the student must obtain at least 20 points out of 25 on Kattis.
Note: The Kattis score does not guarantee the student is given the corresponding grade level, the grade level will depend on later oral examination.

6 Grade level B and A

In order to get B or A, the student is expected to be able to have completed levels E through C and adapt the algorithms where necessary to the checkers problem which can be found on Kattis:

https://kth.kattis.com/problems/kth:ai:checkers.

Having completed levels E through C, it should be a straight-forward to make minimax with alpha-beta pruning work for checkers, however, the student will have to further optimize the algorithm to the problem in order to get the necessary score. To this end the student may pursue several avenues, from iterative deepening to implementing repeated states-checking to move ordering. In other words, it might be possible to get to the required score by fiddling the depth and the evaluation heuristics, but this is not the point of this exercise and will not yield A-B grades.

To achieve this grade the student must obtain at least 20 points out of 25 on Kattis.

Note: The Kattis score does not guarantee the student is given the corresponding grade level, the grade level will depend on later oral examination.